

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels
import statsmodels.api as sm
import sklearn
from sklearn.model_selection import train_test_split
```

```
In [2]: adv=pd.read_csv("https://raw.githubusercontent.com/rkmishracs/dataset/main/advertisi
```

```
In [3]: adv.head()
```

```
Out[3]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [4]: adv.shape
```

```
Out[4]: (200, 4)
```


```
In [5]: adv.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   TV           200 non-null    float64
1   Radio        200 non-null    float64
2   Newspaper    200 non-null    float64
3   Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [6]: adv.describe()
```

Out[6]:


	TV	Radio	Newspaper	Sales
<b>count</b>	200.000000	200.000000	200.000000	200.000000
<b>mean</b>	147.042500	23.264000	30.554000	15.130500
<b>std</b>	85.854236	14.846809	21.778621	5.283892
<b>min</b>	0.700000	0.000000	0.300000	1.600000
<b>25%</b>	74.375000	9.975000	12.750000	11.000000
<b>50%</b>	149.750000	22.900000	25.750000	16.000000
<b>75%</b>	218.825000	36.525000	45.100000	19.050000
<b>max</b>	296.400000	49.600000	114.000000	27.000000

In [7]: `sns.regplot(x='TV', y='Sales', data=adv)`Out[7]: `<AxesSubplot:xlabel='TV', ylabel='Sales'>`In [8]: `sns.regplot(x='Radio', y='Sales', data=adv)`Out[8]: `<AxesSubplot:xlabel='Radio', ylabel='Sales'>`




```
In [9]: sns.regplot(x='Newspaper', y='Sales', data=adv)
```

```
Out[9]: <AxesSubplot:xlabel='Newspaper', ylabel='Sales'>
```



```
In [10]: sns.pairplot(data=adv, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales')
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x241067b26d0>
```




```
In [11]: adv.corr()
```

Out[11]:

	TV	Radio	Newspaper	Sales
TV	1.000000	0.054809	0.056648	0.901208
Radio	0.054809	1.000000	0.354104	0.349631
Newspaper	0.056648	0.354104	1.000000	0.157960
Sales	0.901208	0.349631	0.157960	1.000000


```
In [12]: sns.heatmap(adv.corr())
```

Out[12]: <AxesSubplot:>



```
In [13]: sns.heatmap(adv.corr(), annot=True)
```

Out[13]: <AxesSubplot:>



```
In [14]: X=adv['TV']
         y=adv['Sales']
```

```
In [15]: X_train, X_test, y_train, y_test= train_test_split(X,y, train_size=0.70, random_st
```

```
In [16]: X_train_sm=sm.add_constant(X_train)
         X_train_sm.head()
```

```
Out[16]:
```

	const	TV
74	1.0	213.4
3	1.0	151.5
185	1.0	205.0
26	1.0	142.9
90	1.0	134.3

```
In [17]: lr=sm.OLS(y_train, X_train_sm)
```

```
In [18]: lr_model=lr.fit()
         lr_model.params
```

```
Out[18]: const    6.948683
         TV      0.054546
         dtype: float64
```

```
In [19]: lr_model.summary()
```


Out[19]:

OLS Regression Results							
Dep. Variable:		Sales			R-squared:		0.816
Model:		OLS			Adj. R-squared:		0.814
Method:		Least Squares			F-statistic:		611.2
Date:		Wed, 01 Feb 2023			Prob (F-statistic):		1.52e-52
Time:		11:11:22			Log-Likelihood:		-321.12
No. Observations:		140			AIC:		646.2
Df Residuals:		138			BIC:		652.1
Df Model:		1					
Covariance Type:		nonrobust					
	coef	std err	t	P> t	[0.025	0.975]	
const	6.9487	0.385	18.068	0.000	6.188	7.709	
TV	0.0545	0.002	24.722	0.000	0.050	0.059	
Omnibus:		0.027	Durbin-Watson:		2.196		
Prob(Omnibus):		0.987	Jarque-Bera (JB):		0.150		
Skew:		-0.006	Prob(JB):		0.928		
Kurtosis:		2.840	Cond. No.		328.		


Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [20]: plt.scatter(X_train, y_train)
```

Out[20]: <matplotlib.collections.PathCollection at 0x24107af6e50>



```
In [21]: plt.scatter(X_train, y_train)
plt.plot(X_train, 6.948+0.054*X_train,'r')
plt.show()
```



```
In [22]: y_train_pred=lr_model.predict(X_train_sm)
y_train_pred
```

```
Out[22]: 74      18.588747
          3       15.212365
          185     18.130563
          26      14.743271
          90      14.274178
          ...
          87      12.986898
          103     17.197830
          67      14.546907
          24      10.346884
          8       7.417777
          Length: 140, dtype: float64
```


```
In [23]: residual=y_train-y_train_pred
          residual
```

```
Out[23]: 74      -1.588747
          3       1.287635
          185     4.469437
          26      0.256729
          90     -0.274178
          ...
          87      3.013102
          103     2.502170
          67     -1.146907
          24     -0.646884
          8      -2.617777
          Length: 140, dtype: float64
```


```
In [24]: plt.figure()
          sns.displot(residual)
```

```
Out[24]: <seaborn.axisgrid.FacetGrid at 0x24107bfd5e0>
          <Figure size 640x480 with 0 Axes>
```





```
In [25]: plt.scatter(X_train, residual)
plt.show()
```



```
In [ ]:
```